

# Reinforcement Learning: Core Algorithms

Earl Wong

# Core Algorithms

We examine the following algorithms:

- PPO
- Soft Actor Critic
- DDPG

# PPO

- PPO is a model free, on policy based algorithm (-evolution of the base actor critic algorithm), succeeding TRPO.
- PPO prevents large, positive losses from being back propagated into the network - losses that could destabilize actor / policy training.
- This is achieved by computing the following ratio  $r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)}$  between the current policy and the old policy, and clipping the ratio when it becomes too large.
- Under PPO, the policy improves gradually and steadily.

# PPO

- Paper title: Proximal Policy Optimization Algorithms
- Authors: Schulman, et. al.

# Foundational “Roots”

- PPO can be viewed as the “engineering” implementation / realization of a more theoretically grounded algorithm, TRPO.
- In TRPO, actor / policy updates are performed in a constrained optimization framework.
- A second order conjugate gradient optimizer is used to update the network parameters.
- In TRPO, a KL divergence constraint is applied to the old and current policies, ensuring that the two policies remain close to each other.

# Foundational “Roots”

- PPO “trades off” the guarantees from TRPO, with ease of use and solid empirical results.
- In PPO, actor / policy updates are performed in an unconstrained optimization framework.
- The stochastic gradient descent optimizer is used to update the actor / policy network parameters.
- In PPO, the KL divergence constraint is enforced using a heuristic clipping function.

# PPO

---

**Algorithm 1** PPO, Actor-Critic Style

---

```
for iteration=1, 2, ... do  
  for actor=1, 2, ...,  $N$  do  
    Run policy  $\pi_{\theta_{\text{old}}}$  in environment for  $T$  timesteps  
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$   
  end for  
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$   
   $\theta_{\text{old}} \leftarrow \theta$   
end for
```

---

# The Details

- The first main detail of the algorithm, is the surrogate loss function, used to train the actor / policy.
- In TRPO, the following constraint existed:  $D_{\text{KL}}^{\text{max}}(\theta_{\text{old}}, \theta) \leq \delta$ .
- This constraint required the old policy and current policy to be “close” for every point in the state space.
- Since this constraint was not practical, it was relaxed in the TRPO paper to an “average” KL divergence.



# The Details

- In PPO, clipping is applied to the policy ratio for every point in the state space.
- Specifically  $\min \left( r_t(\theta) \hat{A}_t, \mathbf{clip} \left( r_t(\theta), 1 - \epsilon, 1 + \epsilon \right) \hat{A}_t \right)$
- The clip function discourages back propagating large, positive losses that can destabilize training.
- i.e. It is better to take small, positive steps to encourage actions with positive advantage, rather than taking a single, large, positive step.
- Note: One of the goals of TRPO, was to determine the largest permissible step that could be taken (i.e. trust region) while satisfying a closeness constraint between the old and new policies.
- Note: The surrogate function (shown above) does not protect against taking large, negative steps.

# The Details

- The second main detail involves the advantage estimates  $\hat{A}_t$  that are used to train the critic.
- In practice, generalized advantage estimates are used, to balance the bias - variance tradeoff.
- i.e. If the critic is updated using single time steps (or 1-step advantage estimates), the variance will be low, but the bias will be high.
- i.e. If the critic is updated using the complete monte carlo rollout, the variance will be high, but the bias will be low.

# The Details

- One step advantage estimate:  $\hat{A}_t = r_t + \gamma V(s_{t+1}) - V(s_t)$

- Generalized advantage estimate:  $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$

$$\hat{A}_t = \sum_{l=0}^{T-t-1} (\gamma\lambda)^l \delta_{t+l}$$

- $\lambda$  controls the bias - variance tradeoff.
- If  $\lambda = 1$ , we have the high variance and low bias scenario.
- If  $\lambda \approx 0$ , we have the low variance and high bias scenario.

# The Details

- The final detail involves the introduction of an entropy term in the surrogate loss function, to prevent pre-mature convergence:

$$\mathcal{H}[\pi(\cdot | s_t)] = - \sum_a \pi(a | s_t) \log \pi(a | s_t)$$

- This term helps to keep exploration alive, when the surrogate loss is small.
- In practice, sample based entropy  $\log \pi(a | s_t)$  is used to compute entropy.
- Note: As the confidence in a policy action increases, the entropy will naturally decrease.

# The Details

- PPO can be applied to both continuous and discrete action spaces.
- For discrete action spaces, the logits are converted to action probabilities via the softmax function.
- For continuous action spaces, the network outputs mean and standard deviation values for every input state.
- A gaussian distribution is then created and sampled from, to obtain the continuous action associated with the input state.
- In practice, the reparameterization trick is used, to back propagate through the process to train the network.

# Reparameterization Trick

- The reparameterization trick results from earlier work in variational auto encoders (VAE).
- There, a gaussian distribution was used.
- However, other distributions can also be used.
- i.e. The reparameterization trick can be applied, using other distributions.

# Reparameterization Trick

In practice, gaussian distributions are a “fan favorite” because:

- Gaussian distributions occur naturally in nature.
- Gaussian distributions have nice properties, like smoothness / differentiability.
- Gaussian distributions are mathematically tractable / usually result in closed form, analytic solutions.
- Example: VAE's minimize KL divergence. If both terms are gaussian distributions, a closed form expression results.

# Reparameterization Trick

- Now, what is the trick, and why is it needed?
- First, let's address the need.

Flow:

- For a given state input, the network outputs a mean value (say  $b$ ) and std value (say  $a$ ).
- A normal random variable  $y$  is created, using these output values:

$$y \sim N(b, a^2)$$

- The random variable is then sampled -corresponding to an action.



# Reparameterization Trick

The problem:

- Gradients cannot be back propagated through a sampling process.
- What is the workaround?
- The reparameterization trick.

# Reparameterization Trick

How does it work?

- Define a new function / random variable:  $z = \mu + \sigma \cdot \epsilon$
- Where  $\epsilon \sim N(0,1)$
- By construction,  $z$  is a differentiable function with respect to the mean and standard deviation.
- Now, back propagation can occur.
- Note: Sampling now results from the epsilon random variable.

# Reparameterization Trick

- Does the new random variable  $z$ , have the same distribution as the normal distribution  $y$ , created earlier?
- Yes.
- Why?
- We have applied a linear transformation  $aX + b$ , to a  $N(0,1)$  random variable with mean  $b$  and standard deviation  $a$ .
- The resulting random variable  $z$ , has the same distribution as

$$y \sim N(b, a^2)$$

# Soft Actor Critic

- Soft actor critic is a model free, off policy algorithm (-evolution of the base actor critic algorithm), designed for continuous action spaces.
- Soft actor critic employs a stochastic policy - usually gaussian.
- Soft actor critic introduces entropy directly into the objective function (vs the loss function).
- As a result, the objective changes from reward maximization to BOTH reward maximization and entropy maximization.
- Soft actor critic employs two critics to reduce over estimation biases.

# Soft Actor Critic

- Paper title: Soft Actor-Critic: Off Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor
- Authors: Haarnoja, Zhou, Abbeel, Levine

# Foundational “Roots”

- In PPO, entropy was introduced into the LOSS FUNCTION, to keep training alive / prevent early convergence to a local minima.
- In SAC, entropy is directly infused into the TRAINING OBJECTIVE, to encourage BOTH high reward and high exploration.
- In SAC, entropy is also introduced into the two target Q networks, used to train the two main Q networks (critics).
- The latter entropy helps keep training alive / prevents early convergence of the two main Q networks (critics), to a local minima.

# Foundational “Roots”

- In SAC, continuous actions are obtained by sampling a gaussian distribution.
- The gaussian distribution is created from  $\mu$  and  $\sigma$  vector outputs from a neural network, for a given state vector input.
- In order to successfully perform back propagation through this structure, the reparameterization trick is employed.
- The reparameterization trick was first introduced during the training of variational auto encoders.

# Foundational “Roots”

- Whenever  $\max_a Q(s, a)$  appears in a policy, over estimation biases result.
- This is because of the inherent noise present, in the action value function output.
- i.e. By taking the max of noisy expected returns, over estimation results.
- Over estimation bias was first addressed by Van Hasselt in 2010 for the tabular data case, with the introduction of two Q-tables.
- Soft Actor Critic addresses overestimation bias by 1) estimating two Q networks (two Q critics) and 2) taking the min of the values produced.



# Foundational “Roots”

- Like DQN, training stability of the Q network was of paramount importance.
- With DQN, a target network was introduced, storing a “delayed” version of the main Q network parameters, to enhance learning stability.
- In Soft Actor Critic, two target Q networks are introduced alongside the two main Q networks (critics).
- In Soft Actor Critic, polyak averaging is used to update the parameters of the two target Q networks.

# Foundational “Roots”

- Since SAC is an offline algorithm, SAC samples from and writes to (=updates the inputs in) a replay buffer.
- This activity parallels the activity described in DQN algorithm training.
- i.e. Experience replay and sampling from a replay buffer.

# Soft Actor Critic

## - original paper

---

**Algorithm 1** Soft Actor-Critic

---

Initialize parameter vectors  $\psi, \bar{\psi}, \theta, \phi$ .  
**for** each iteration **do**  
    **for** each environment step **do**  
         $\mathbf{a}_t \sim \pi_\phi(\mathbf{a}_t|\mathbf{s}_t)$   
         $\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1}|\mathbf{s}_t, \mathbf{a}_t)$   
         $\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$   
    **end for**  
    **for** each gradient step **do**  
         $\psi \leftarrow \psi - \lambda_V \hat{\nabla}_\psi J_V(\psi)$   
         $\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i)$  for  $i \in \{1, 2\}$   
         $\phi \leftarrow \phi - \lambda_\pi \hat{\nabla}_\phi J_\pi(\phi)$   
         $\bar{\psi} \leftarrow \tau\psi + (1 - \tau)\bar{\psi}$   
    **end for**  
**end for**

---

# Soft Actor Critic - from openAI

## Pseudocode

---

**Algorithm 1** Soft Actor-Critic

---

- 1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi_1, \phi_2$ , empty replay buffer  $\mathcal{D}$
- 2: Set target parameters equal to main parameters  $\phi_{\text{targ},1} \leftarrow \phi_1, \phi_{\text{targ},2} \leftarrow \phi_2$
- 3: **repeat**
- 4:   Observe state  $s$  and select action  $a \sim \pi_\theta(\cdot|s)$
- 5:   Execute  $a$  in the environment
- 6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
- 7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$
- 8:   If  $s'$  is terminal, reset environment state.
- 9:   **if** it's time to update **then**
- 10:     **for**  $j$  in range(however many updates) **do**
- 11:       Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$
- 12:       Compute targets for the Q functions:

$$y(r, s', d) = r + \gamma(1 - d) \left( \min_{i=1,2} Q_{\phi_{\text{targ},i}}(s', \tilde{a}') - \alpha \log \pi_\theta(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_\theta(\cdot|s')$$

- 13:     Update Q-functions by one step of gradient descent using

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2 \quad \text{for } i = 1, 2$$

- 14:     Update policy by one step of gradient ascent using

$$\nabla_\theta \frac{1}{|B|} \sum_{s \in B} \left( \min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_\theta(s)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s)|s) \right),$$

where  $\tilde{a}_\theta(s)$  is a sample from  $\pi_\theta(\cdot|s)$  which is differentiable wrt  $\theta$  via the reparametrization trick.

- 15:     Update target networks with

$$\phi_{\text{targ},i} \leftarrow \rho \phi_{\text{targ},i} + (1 - \rho) \phi_i \quad \text{for } i = 1, 2$$

- 16:     **end for**
- 17:   **end if**
- 18: **until** convergence

# The Details

- Soft Actor Critic maximizes both return and exploration.

$$\nabla_{\theta} \frac{1}{|B|} \sum_{s \in B} \left( \min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_{\theta}(s)) - \alpha \log \pi_{\theta}(\tilde{a}_{\theta}(s) | s) \right),$$

- Soft Actor Critic trains two main Q networks (critics), using the squared error loss.

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s,a) - y(r,s',d))^2$$

# The Details

- In the first equation, entropy is specifically added to the training objective.
- Recall (previous):

$$\pi_* = \operatorname{argmax}_{\pi} E_{\tau \sim \pi} \left[ \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \right]$$

- Now:

$$\pi_* = \operatorname{argmax}_{\pi} E_{\tau \sim \pi} \left[ \sum_{t=0}^{\infty} \gamma^t (R(s_t, a_t) + \alpha H(\pi(\cdot | s_t))) \right]$$

# The Details

- The state value function now becomes the “soft” version:

$$V_{\pi}(s) = E_{\tau \sim \pi} \left[ \sum_{t=0}^{\infty} \gamma^t (R(s_t, a_t) + \alpha H(\pi(\cdot | s_t))) \mid s_0 = s \right]$$

- The action value function now becomes the “soft” version:

$$Q_{\pi}(s, a) = \mathbb{E}_{\tau \sim \pi} \left[ \sum_{t=0}^{\infty} \gamma^t \left( R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot | s_t)) \right) \mid s_0 = s, a_0 = a \right]$$

- Linking state and action value functions via entropy:

$$V_{\pi}(s) = E_{a \sim \pi} [Q_{\pi}(s, a)] + \alpha H(\pi(\cdot | s))$$

# The Details

- With the addition of entropy, the Bellman Equation for the Q network now becomes:

$$Q_{\pi}(s, a) = \mathbb{E}_{\substack{s' \sim P(\cdot | s, a) \\ a' \sim \pi(\cdot | s')}} \left[ R(s, a, s') + \gamma \left( Q_{\pi}(s', a') + \alpha H(\pi(\cdot | s')) \right) \right]$$

- Or:

$$Q_{\pi}(s, a) = \mathbb{E}_{\substack{s' \sim P(\cdot | s, a) \\ a' \sim \pi(\cdot | s')}} \left[ R(s, a, s') + \gamma V_{\pi}(s') \right]$$



# The Details

- Soft Actor Critic samples from a gaussian distribution, and uses the reparameterization trick to back propagate through the network.
- In python code:

```
def sample_action_manual(mu, std):  
    epsilon = torch.randn_like(std)           # Sample from N(0, 1)  
    action = mu + std * epsilon               # Reparameterized sample  
    return action
```

```
def sample_action_rsample(mu, std):  
    dist = torch.distributions.Normal(mu, std)  
    action = dist.rsample()  
    return action
```

```
mu, std = policy(state)  
  
# Manual  
action_manual = sample_action_manual(mu, std)  
  
# One-liner  
action_rsample = sample_action_rsample(mu, std)
```

# The Details

- Soft Actor Critic reduces over estimation bias by utilizing two target Q networks, and selecting the action that results in the minimum return.

$$y(r, s', d) = r + \gamma(1 - d) \left( \min_{i=1,2} Q_{\phi_{\text{targ},i}}(s', \tilde{a}') - \alpha \log \pi_{\theta}(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_{\theta}(\cdot|s')$$

- Soft Actor Critic introduces entropy into the target network used for main Q network training, to reduce the chance of early convergence to a local minima.
- Soft Actor Critic updates the target network using polyak averaging between the main and target network parameters.

$$\phi_{\text{targ},i} \leftarrow \rho \phi_{\text{targ},i} + (1 - \rho) \phi_i$$

# DDPG

- Deep deterministic Policy Gradient (DDPG) is a model free, off policy based actor-critic algorithm.
- DDPG extends DQN, by producing a policy with continuous actions, enabling the algorithm to be applied to physical control tasks.
- DDPG learns policies with deterministic actions.

# DDPG

- Paper title: Continuous Control with Deep Reinforcement Learning
- Authors: Lillicrap, et. al.

# Foundational “Roots”

- DDPG extends DQN to continuous action spaces.
- DDPG uses the same two tricks to train DQN: 1) target network and 2) replay buffer with temporally uncorrelated / near IID inputs.
- DDPG applies normalization across input features, much like classical machine learning approaches.

# DDPG

---

**Algorithm 1** DDPG algorithm

---

Randomly initialize critic network  $Q(s, a|\theta^Q)$  and actor  $\mu(s|\theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .

Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer  $R$

**for** episode = 1, M **do**

    Initialize a random process  $\mathcal{N}$  for action exploration

    Receive initial observation state  $s_1$

**for**  $t = 1, T$  **do**

        Select action  $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise

        Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$

        Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$

        Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$

        Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'}))|\theta^{Q'}$

        Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

        Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

    Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

**end for**

**end for**

---

# The Details

- Like REINFORCE in vanilla policy gradient, DDPG wants to find the network weights that will maximize the return.

$$J(\theta) = E_{\tau \sim u} \left[ \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right]$$

- However, instead of maximizing the return over a trajectory using a stochastic policy, the return is now maximized over a deterministic policy based on a state visitation distribution.

$$J(\theta) = E_{s \sim \rho^u} [Q^u(s, u_{\theta}(s))]$$

- $\rho^u(s)$  Is used to denote the state visitation distribution.

# The Details

- It is formally defined as:

$$\rho^u(s) = \sum_{t=0}^{\infty} \gamma^t P(s_t = s | u_\theta)$$

- The objective function can now be written as:

$$J(\theta) = \int \rho^u(s) r(s, u_\theta(s)) ds = E_{s \sim \rho^u} [Q_u(s, u_\theta(s))]$$

- Differentiating the objective, we obtain:

$$\nabla_\theta J(\theta) = \int \rho^u(s) \nabla_\theta Q_u(s, u_\theta(s)) ds$$

$$\nabla_\theta J(\theta) = \int \rho^u(s) \nabla_a Q_u(s, a) |_{a=u_\theta(s)} \nabla_\theta u_\theta(s) ds$$

$$\nabla_\theta J(\theta) = E_{s \sim \rho^u} \left[ \nabla_a Q_u(s, a) |_{a=u_\theta(s)} \nabla_\theta u_\theta(s) \right]$$



# The Details

- Note:  $\rho^u(s)$  does depend on the network parameters.
- However, the author has chosen to ignore this “indirect” dependency, and the corresponding chain rule term associated with it.

# The Details

- DDPG encourages exploration by adding a noise term to the deterministic action produced by the policy.

$$a_t = u(s_t | \theta^u) + N_t$$

# The Details

- DDPG uses two target networks -one for the Q function and one for the target policy:

$$\theta^{Q'} \leftarrow \theta^Q, \theta^{u'} \leftarrow \theta^u$$

- DDPG updates the network parameters in the target Q function and target policy using weighted averaging:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{u'} \leftarrow \tau \theta^u + (1 - \tau) \theta^{u'}$$