

# Reinforcement Learning: Visual Control

Earl Wong

# Visual Control Algorithms

We examine the following algorithms:

- GuidedPolicy Search
- DQN
- DrQ-v2
- Dreamer-v3

# DQN

- DQN learned successful policies from high dimensional inputs using end to end reinforcement learning.
- DQN used a deep Q-network, to output actions from pixel based inputs.
- DQN training made use of 1) replay buffers / experience replay and 2) an additional target Q network.
- Paper title: Human-level control through deep reinforcement learning
- Authors: Mnih, et. al.

# Foundational “Roots”

- DQN can be viewed as classical (tabular) Q learning, replaced by a deep neural network.
- Like classical Q learning, DQN employed two critics, to address over estimation biases.

# DQN

**Algorithm 1: deep Q-learning with experience replay.**

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value function  $Q$  with random weights  $\theta$

Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$

**For** episode = 1,  $M$  **do**

    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$

**For**  $t = 1, T$  **do**

        With probability  $\varepsilon$  select a random action  $a_t$

        otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$

        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$

        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$

        Every  $C$  steps reset  $\hat{Q} = Q$

**End For**

**End For**

# The Details

- DQN made use of two action value functions - the Q function associated with the main network:  $Q$ , and the Q function associated with the target network:  $\hat{Q}$ .
- $\hat{Q}$  was used to compute a target  $y = r_j + \gamma \cdot \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-)$
- The difference between the target  $y$  and  $Q$ :  $(y_j - Q(\phi_j, a_j; \theta))^2$  was then back propagated into main Q network at each iteration.
- After a finite number of iterations / back propagations, the weights from  $Q$  were (then) copied over the existing weights associated with  $\hat{Q}$ .
- The process then repeated.

# The Details

- This training “trick” was instrumental in increasing the training stability of the deep neural network.
- Today, a variation of this trick is used.
- Instead of performing an update (via copy) after a finite number of iterations, the update now occurs at every iteration using a weighted average (polyak averaging).

# The Details

- The DQN algorithm sampled from a buffer, consisting of input vectors:  $(\phi_j, a_j, r_j, \phi_{j+1})$
- $\phi_j$  consisted of 4 temporally correlated video frames.
- After an action was taken, the environment then output a new video frame  $x_{t+1}$ .
- $\phi_j$  was then updated from  $\phi_j = \{x_{t-3}, x_{t-2}, x_{t-1}, x_t\}$  to  $\phi_{j+1} = \{x_{t-2}, x_{t-1}, x_t, x_{t+1}\}$
- Valuable temporal information - such as velocity, was contained in the temporally correlated frames.



# The Details

- A mini batch is created, by sampling from a replay buffer containing the (previously described) vectors.
- Although the four input frames  $\phi_j$  in each vector are temporally correlated, the vectors in the mini batch are not.
- This is because the vectors are sampled from different “plays” of the same game and / or different time instances.
- i.e. The vectors in the mini batch are approximately IID.
- DQN training was enhanced, by using mini batches with non-temporally correlated vectors.

# DrQ-v2

- DrQ-v2 is a model free, off policy, actor-critic RL algorithm, for continuous visual control.
- It is the first model free RL algorithm to solve humanoid locomotion (stand, walk, run) using pixel based data.
- Paper title: Mastering Continuous Visual Control: Improved Data Augmented Reinforcement Learning
- Authors: Yarats, Fergus, Lazaric, Pinto

# Foundational “Roots”

- DrQ-v2 uses data augmentation from computer vision, to achieve state of the art results for a suite of Deep Mind Control (DMC) tasks.
- DrQ-v2 uses the DDPG algorithm for it's actor-critic implementation (vs SAC for DrQ) because it allowed n-step returns to be easily computed.
- DrQ-v2 uses entropy, to keep learning alive.
- DrQ-v2 uses two main Q networks (to reduce over estimation biases) and two target Q networks (to facilitate learning).

# DrQ-v2

---

**Algorithm 1** DrQ-v2: Improved data-augmented RL.

---

**Inputs:**

$f_\xi, \pi_\phi, Q_{\theta_1}, Q_{\theta_2}$ : parametric networks for encoder, policy, and Q-functions respectively.

aug: random shifts image augmentation.

$\sigma(t)$ : scheduled standard deviation for the exploration noise defined in Equation (3).

$T, B, \alpha, \tau, c$ : training steps, mini-batch size, learning rate, target update rate, clip value.

**Training routine:**

**for** each timestep  $t = 1..T$  **do**

$\sigma_t \leftarrow \sigma(t)$

$\mathbf{a}_t \leftarrow \pi_\phi(f_\xi(\mathbf{x}_t)) + \epsilon$  and  $\epsilon \sim \mathcal{N}(0, \sigma_t^2)$

$\mathbf{x}_{t+1} \sim P(\cdot | \mathbf{x}_t, \mathbf{a}_t)$

$\mathcal{D} \leftarrow \mathcal{D} \cup (\mathbf{x}_t, \mathbf{a}_t, R(\mathbf{x}_t, \mathbf{a}_t), \mathbf{x}_{t+1})$

    UPDATECRITIC( $\mathcal{D}, \sigma_t$ )

    UPDATEACTOR( $\mathcal{D}, \sigma_t$ )

**end for**

**procedure** UPDATECRITIC( $\mathcal{D}, \sigma$ )

$\{(\mathbf{x}_t, \mathbf{a}_t, r_{t:t+n-1}, \mathbf{x}_{t+n})\} \sim \mathcal{D}$

$\mathbf{h}_t, \mathbf{h}_{t+n} \leftarrow f_\xi(\text{aug}(\mathbf{x}_t)), f_\xi(\text{aug}(\mathbf{x}_{t+n}))$

$\mathbf{a}_{t+n} \leftarrow \pi_\phi(\mathbf{h}_{t+n}) + \epsilon$  and  $\epsilon \sim \text{clip}(\mathcal{N}(0, \sigma^2))$

    Compute  $\mathcal{L}_{\theta_1, \xi}$  and  $\mathcal{L}_{\theta_2, \xi}$  using Equation (1)

$\xi \leftarrow \xi - \alpha \nabla_\xi (\mathcal{L}_{\theta_1, \xi} + \mathcal{L}_{\theta_2, \xi})$

$\theta_k \leftarrow \theta_k - \alpha \nabla_{\theta_k} \mathcal{L}_{\theta_k, \xi} \quad \forall k \in \{1, 2\}$

$\bar{\theta}_k \leftarrow (1 - \tau)\bar{\theta}_k + \tau\theta_k \quad \forall k \in \{1, 2\}$

**end procedure**

**procedure** UPDATEACTOR( $\mathcal{D}, \sigma$ )

$\{(\mathbf{x}_t)\} \sim \mathcal{D}$

$\mathbf{h}_t \leftarrow f_\xi(\text{aug}(\mathbf{x}_t))$

$\mathbf{a}_t \leftarrow \pi_\phi(\mathbf{h}_t) + \epsilon$  and  $\epsilon \sim \text{clip}(\mathcal{N}(0, \sigma^2))$

    Compute  $\mathcal{L}_\phi$  using Equation (2)

$\phi \leftarrow \phi - \alpha \nabla_\phi \mathcal{L}_\phi$

**end procedure**

---

▷ Compute stddev for the exploration noise

▷ Add noise to the deterministic action

▷ Run transition function for one step

▷ Add a transition to the replay buffer

▷ Sample a mini batch of  $B$  transitions

▷ Apply data augmentation and encode

▷ Sample action

▷ Compute critic losses

▷ Update encoder weights

▷ Update critic weights

▷ Update critic target weights

▷ Sample a mini batch of  $B$  observations

▷ Apply data augmentation and encode

▷ Sample action

▷ Compute actor loss

▷ Update actor's weights only

# The Details

- $h_t \leftarrow f_{\xi}(aug(x_t))$  takes an input image  $x_t$  of size 84x84, pads it (using nearest neighbor replication of value 4) and then randomly samples new 84x84 crops from it.
- Each pixel in the crop is blurred, using the average of its four nearest neighbors.
- The blurred / averaged image is then encoded (via neural network) into a low dimensional latent space vector  $h_t$

# The Details

- DrQ-v2 uses N step returns, to improve training efficiency.
- To accomplish this, DrQ-v2 changed from a soft actor critic (SAC) backbone to a deep deterministic policy gradient (DDPG) backbone.
- In DDPG, exploration was no longer built into the objective function (like in SAC), but added as a noise term.

# The Details

- DrQ-v2 added a “twist” to the original DPPG exploration construct, by making the noise time dependent.
- i.e.  $a_t \leftarrow \pi_\phi(f_\xi(x_t)) + \epsilon$  where  $\epsilon \sim N(0, \sigma^2(t))$
- Now, exploration was based on a time dependent schedule.

# The Details

- DrQ-v2 was computationally efficient and computationally competitive, because it was re-engineered to have a fast replay buffer and to perform fast data augmentation.
- DrQ-v2 performed well on hard DMC tasks (=tasks with large initial state distributions) by significantly increasing the size of the replay buffer.



# Dreamer-v3

- Dreamer-v3 is a model based reinforcement learning (MBRL) algorithm.
- Dreamer-v3 learns a policy using an actor-critic algorithm trained on “latent imagination” produced by the Dreamer world model.
- Dreamer-v3 is a “generalist” algorithm, that was successfully trained on over 150 diverse visual tasks (Atari, ProcGen, DMLab, VisualControl, BSuite), outperforming many specialized models.
- Dreamer-v3 achieved this using 1) a fixed architecture, 2) fixed loss functions and 3) fixed hyper parameter settings.

# Dreamer-v3

- Paper title: Mastering Diverse Control Tasks Through World Models
- Authors: Hafner, Pasukonis, Ba, Lillicrap


# Foundational “Roots”

- The learned world model consisted of a recurrent state space model (RSSM) / latent dynamics model, a rewards predictor and a reconstruction predictor.
- Instead of interacting with the environment or using buffered data / experience replay, a policy was now learned using the world model environment / outputs from the world model.
- i.e. Learning a policy from rollouts / imagined trajectories in the latent space, where the world model supplies the next latent states and rewards.

# Foundational “Roots”

- Dreamer-v3 used KL divergence to ensure that the world model latent state outputs matched the latent state outputs computed from real visual information.
- Dreamer-v3 was able to learn policies for 150 diverse tasks (=good generalization) because it performed variable normalization.
- Specifically, normalization was applied to the input frames, the predicted output frames, the predicted rewards and the computed value function.

# Dreamer-v3



- 1) The Dreamer world model was realized using a recurrent state space model (RSSM), a rewards predictor and a reconstruction predictor.
- 2) An actor-critic algorithm was trained to determine the best actions, using future latent states and rewards produced by the world model.

# Dreamer-v1 (Original)

---

**Algorithm 1:** Dreamer

---

Initialize dataset  $\mathcal{D}$  with  $S$  random seed episodes.  
Initialize neural network parameters  $\theta, \phi, \psi$  randomly.  
**while not converged do**  
    **for update step  $c = 1..C$  do**  
        // Dynamics learning  
        Draw  $B$  data sequences  $\{(a_t, o_t, r_t)\}_{t=k}^{k+L} \sim \mathcal{D}$ .  
        Compute model states  $s_t \sim p_\theta(s_t | s_{t-1}, a_{t-1}, o_t)$ .  
        Update  $\theta$  using representation learning.  
        // Behavior learning  
        Imagine trajectories  $\{(s_\tau, a_\tau)\}_{\tau=t}^{t+H}$  from each  $s_t$ .  
        Predict rewards  $E(q_\theta(r_\tau | s_\tau))$  and values  $v_\psi(s_\tau)$ .  
        Compute value estimates  $V_\lambda(s_\tau)$  via [Equation 6](#).  
        Update  $\phi \leftarrow \phi + \alpha \nabla_\phi \sum_{\tau=t}^{t+H} V_\lambda(s_\tau)$ .  
        Update  $\psi \leftarrow \psi - \alpha \nabla_\psi \sum_{\tau=t}^{t+H} \frac{1}{2} \|v_\psi(s_\tau) - V_\lambda(s_\tau)\|^2$ .  
    // Environment interaction  
     $o_1 \leftarrow \text{env.reset}()$   
    **for time step  $t = 1..T$  do**  
        Compute  $s_t \sim p_\theta(s_t | s_{t-1}, a_{t-1}, o_t)$  from history.  
        Compute  $a_t \sim q_\phi(a_t | s_t)$  with the action model.  
        Add exploration noise to action.  
         $r_t, o_{t+1} \leftarrow \text{env.step}(a_t)$ .  
    Add experience to dataset  $\mathcal{D} \leftarrow \mathcal{D} \cup \{(o_t, a_t, r_t)_{t=1}^T\}$ .

---

**Model components**

Representation	$p_\theta(s_t   s_{t-1}, a_{t-1}, o_t)$
Transition	$q_\theta(s_t   s_{t-1}, a_{t-1})$
Reward	$q_\theta(r_t   s_t)$
Action	$q_\phi(a_t   s_t)$
Value	$v_\psi(s_t)$

**Hyper parameters**

Seed episodes	$S$
Collect interval	$C$
Batch size	$B$
Sequence length	$L$
Imagination horizon	$H$
Learning rate	$\alpha$

# Evolution

- Dreamer-v2 and Dreamer-v3 did not provide a “cut and paste” algorithm descriptions.
- This was (probably) attributed to the “evolution details” for each new paper.
- However, the underlying methodology remained constant.
- 1) Use real data (image, action, reward) to train a world model.
- 2) Freeze the world model, and then train the actor-critic model using imagined, latent space trajectories.

# Evolution

- Discrete action outputs were supported in v2.
- Loss functions definitions were modified between v1, v2 and v3.
- Normalizations were applied in Dreamer-v3, enabling Dreamer-v3 to train on 150 tasks without the need for retuning, etc.



# The Details

- The Dreamer-v3 algorithm has the networks: the world model, the actor and the critic.
- The world model consisted of the RSSM (sequence model, encoder and dynamics predictor), the reward predictor, the decoder and the continue predictor.

Sequence model:	$h_t = f_\phi(h_{t-1}, z_{t-1}, a_{t-1})$
Encoder:	$z_t \sim q_\phi(z_t   h_t, x_t)$
Dynamics predictor:	$\hat{z}_t \sim p_\phi(\hat{z}_t   h_t)$
Reward predictor:	$\hat{r}_t \sim p_\phi(\hat{r}_t   h_t, z_t)$
Continue predictor:	$\hat{c}_t \sim p_\phi(\hat{c}_t   h_t, z_t)$
Decoder:	$\hat{x}_t \sim p_\phi(\hat{x}_t   h_t, z_t)$

- All of the components in the world model were trained concurrently.

# The Details

- The world model loss had 3 different loss terms:

$$\mathcal{L}(\phi) \doteq E_{q_\phi} \left[ \sum_{t=1}^T (\beta_{\text{pred}} \mathcal{L}_{\text{pred}}(\phi) + \beta_{\text{dyn}} \mathcal{L}_{\text{dyn}}(\phi) + \beta_{\text{rep}} \mathcal{L}_{\text{rep}}(\phi)) \right]$$

$$\begin{aligned} \mathcal{L}_{\text{pred}}(\phi) &\doteq -\log p_\phi(x_t|z_t, h_t) - \log p_\phi(r_t|z_t, h_t) - \log p_\phi(c_t|z_t, h_t) \\ \mathcal{L}_{\text{dyn}}(\phi) &\doteq \max(1, \text{KL}[\text{sg}(q_\phi(z_t|h_t, x_t)) \| p_\phi(z_t|h_t)]) \\ \mathcal{L}_{\text{rep}}(\phi) &\doteq \max(1, \text{KL}[q_\phi(z_t|h_t, x_t) \| \text{sg}(p_\phi(z_t|h_t))]) \end{aligned}$$

- pred lumped together the reward, continue and decoder loss.
- dyn handled the difference between the predicted latent space output for the world model and the latent space output trained using input data.
- rep should be for the sequence model? (check code)

# The Details

- The symlog function was applied to the encoder inputs and the reconstruction outputs.
- The symlog function compressed large positive and negative values, while simultaneously preserving symmetry.
- The network was trained to produce compressed, reconstruction outputs.
- symexp was then applied, to produce the true reconstruction outputs.

$$\mathcal{L}(\theta) \doteq \frac{1}{2} (f(x, \theta) - \text{symlog}(y))^2 \quad \hat{y} \doteq \text{symexp}(f(x, \theta))$$

$$\begin{aligned} \text{symlog}(x) &\doteq \text{sign}(x) \log(|x| + 1) \\ \text{symexp}(x) &\doteq \text{sign}(x) (\exp(|x|) - 1) \end{aligned}$$

# The Details

- The reward and value function scalers were computed using 2 hot encoding.
- First, a compressed, output range was determined - [min, max].
- Then, a fixed number of histogram bins were used to subdivide this range.
- i.e. Each bin was assigned a specific value in the range.

# The Details

- symlog was then applied, to compress the scalar values to lie within the  $[\min, \max]$  range.
- By construction, each compressed scalar value would lie between two consecutive bins.
- These adjacent bins were then assigned probabilities (summing to 1), based on their closeness to the scalar value.

# The Details

- At the same time, the network produced a sequence of logit values associated with each bin location.
- 2 hot cross entropy was then applied, to compute the resulting loss.

$$\mathcal{L}(\theta) \doteq - \text{twohot}(y)^T \log \text{softmax}(f(x, \theta))$$

$$\hat{y} \doteq \text{softmax}(f(x))^T B \quad B \doteq \text{symexp}([-20 \quad \dots \quad +20])$$

- The authors claimed that this was better for propagating gradients with large targets with large variance.
- i.e. The process was more stable than traditional regression.

# Example

- The range is  $[-20, 20]$ .
- 9 bins are assigned to this range.
- Following symbol compression, the scalar value maps to -19.
- This scalar lies between bin0 and bin1, with assigned values of -20 and -16.
- As a result, bin0 receives a probability of 0.75 while bin1 receives a probability of 0.25.

# Example

- The network outputs 9 logits for the predicted scalar quantity.
- These logits are reprocessed using the softmax function.
- Suppose the result is [.20, .15, .01, .30, .11, .20, .01, .01, .01]
- 2 hot cross entropy loss is then performed:  $.75 \cdot .20 + .25 \cdot .15$
- This loss is then back propagated into the network.